

EAI Diffusion Reading Group

Poisson Field Generative Models

01/04/2023

Summary

1. The Poisson Equation
2. PFGM Derivation
3. Results
4. PFGM++

The Poisson Equation

Given some *source density* $\rho : \mathbb{R}^N \rightarrow \mathbb{R}$ (with *compact support* *)

$$\Delta\varphi = -\rho$$

- $\Delta = \sum_{\mu=1}^N \frac{\partial^2}{\partial x_\mu^2}$ Laplacian operator (again!)
- Poisson Field $E(x) = -\nabla\varphi(x)$ satisfies *Gauss' Law*

$$\nabla \cdot E = \rho$$

One of the most basic PDEs in physics (yields Newtonian gravity, Electrostatic theory, ...)

* ρ is zero outside of some bounded set

Analytical solution

Given by *Green's function*

$$\varphi(x) = \int G(x, y) \rho(y) dy$$

- $G(x, y) = \frac{1}{(N-2)S_{N-1}} \frac{1}{\|x-y\|^{N-2}}$ (Green's function: solution for a single point source)

- $$E(x) = - \int \nabla_x G(x, y) \rho(y) dy$$

- $$\nabla_x G(x, y) = \frac{1}{S_{N-1}} \frac{x-y}{\|x-y\|^N}$$

Gradient Flow


Consider a charged particle dropped in the field E , its trajectory is given by

$$\frac{dx}{dt} = E(x)$$

For a distribution of particles:

$$\frac{\partial p_t}{\partial t}(x) = -\nabla \cdot (p_t(x)E(x))$$

Field Effects



Finally putting the "consciousness" in large Neural Nets

Rescalable dynamics

The Gradient flow ODE is stiff :(. Luckily, for $f \in \mathcal{C}^1$, strictly positive

$$\frac{dx}{dt} = \pm f(x)E(x)$$

follows the same trajectories as the original gradient flow

Forward Process

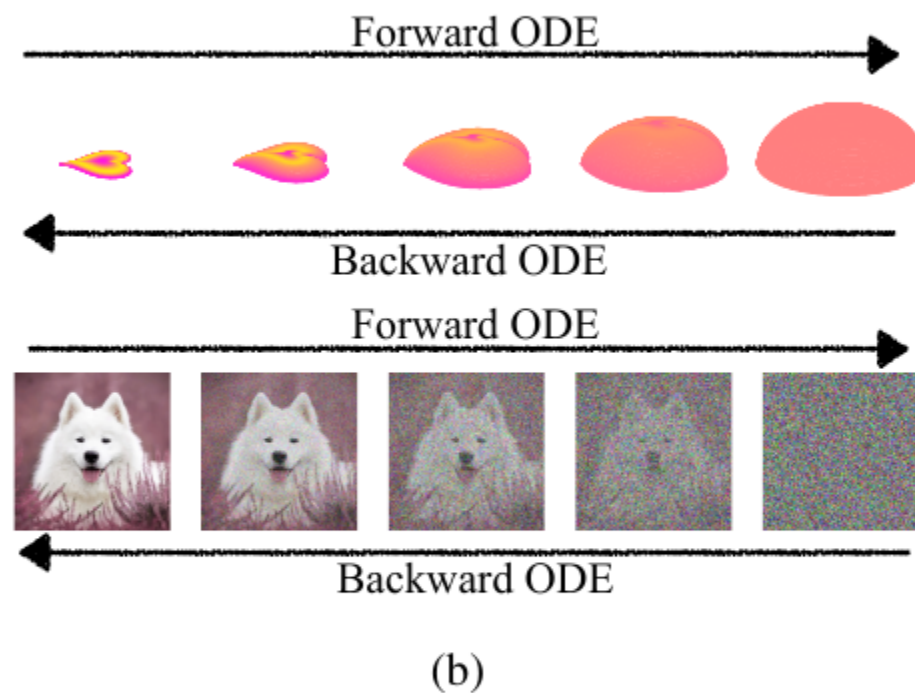



Figure 1: **(a)** 3D Poisson field trajectories for a heart-shaped distribution **(b)** The evolvments of a distribution (**top**) or an (augmented) sample (**bottom**) by the forward/backward ODEs pertained to the Poisson field.


From the Physics to the actual forward process

Several steps needed


1. Add extra dimension z (avoid mode collapse)
2. Normalize field E (numerical stability)
3. Replace time with z (for easier batching)

Add extra dimension


No augmentation (2D)



Augmentation (3D)



(a)



(b)

Figure 2: **(a)** Poisson field (black arrows) and particle trajectories (blue lines) of a 2D uniform disk (red). **Left** (no augmentation, 2D): all particles collapse to the disk center. **Right** (augmentation, 3D): particles hit different points on the disk. **(b)** Proof idea of Theorem 1. By Gauss's Law, the outflow flux $d\Phi_{out}$ equals the inflow flux $d\Phi_{in}$. The factor of two in $p(\mathbf{x})dA/2$ is due to the symmetry of Poisson fields in $z < 0$ and $z > 0$.

Add extra dimension

- Instead of solving the Poisson equation $\Delta\varphi = -p$ in \mathbb{R}^N , solve in augmented space $\tilde{x} = (x, z) \in \mathbb{R}^{N+1}$
- Embed original data with $x \mapsto (x, 0)$

Asymptotics as $\|x\| \rightarrow \infty$

Since p has bounded support, as x follows $\frac{dx}{dt} = E(x)$, E becomes spherically symmetric

This means that for a sufficiently large radius r , the distribution p_t is effectively uniform on the hemi-sphere of radius r

Normalize E

Given a dataset $\mathcal{D} = \{x_i\}_{i=1}^n$, the empirical Poisson field is

$$\hat{E}(\tilde{x}) = c(\tilde{x}) \sum_{i=1}^n \frac{\tilde{x} - \tilde{x}_i}{\|\tilde{x} - \tilde{x}_i\|^{N+1}}$$

where $c(\tilde{x}) = 1 / \sum_{i=1}^n \frac{1}{\|\tilde{x} - \tilde{x}_i\|^{N+1}}$

Further normalization yields the *negative normalized field*

$$v(\tilde{x}) = -\frac{\sqrt{N}\hat{E}(\tilde{x})}{\|\hat{E}(\tilde{x})\| + \gamma}$$

Notation: fields calculated from batch \mathcal{B} : $\hat{E}_{\mathcal{B}}$, $v_{\mathcal{B}}$

NB: γ : small constant to avoid division by zero

Perturbing the augmented training data

Given training datapoint $x \in \mathbb{R}^N$, add noise to $\tilde{x} = (x, 0)$ to obtain (y, z) :

$$y = x + \|\epsilon_x\|(1 + \tau)^m u, \quad z = |\epsilon_z|(1 + \tau)^m$$

- $\epsilon = (\epsilon_x, \epsilon_z) \sim \mathcal{N}(0, \sigma^2 I)$
- $u \sim \mathcal{U}(S_N)$
- $m \sim \mathcal{U}[0, M]$

Hyperparameters: τ, σ, M

Loss

Given mini-batch data $\mathcal{B} = \{x_i\}_{i=1}^{|\mathcal{B}|}$ and a larger batch \mathcal{B}_L (for estimating the normalized field), we train a network f_θ to minimize

$$\mathcal{L}(\theta) = \frac{1}{|\mathcal{B}|} \sum_{i=1}^{|\mathcal{B}|} \|f_\theta(\tilde{y}_i) - v_{\mathcal{B}_L}(\tilde{y}_i)\|_2^2$$

Training Algorithms

Algorithm 1 Learning the normalized Poisson Field

Input: Training iteration T , Initial model f_θ , dataset \mathcal{D} , constant γ , learning rate η .

for $t = 1 \dots T$ **do**

 Sample a large batch \mathcal{B}_L from \mathcal{D} and subsample a batch of datapoints $\mathcal{B} = \{\mathbf{x}_i\}_{i=1}^{|\mathcal{B}|}$ from \mathcal{B}_L

 Simulate the ODE: $\{\tilde{\mathbf{y}}_i = \text{perturb}(\mathbf{x}_i)\}_{i=1}^{|\mathcal{B}|}$

 Calculate the normalized field by \mathcal{B}_L : $\mathbf{v}_{\mathcal{B}_L}(\tilde{\mathbf{y}}_i) = -\sqrt{N}\hat{\mathbf{E}}_{\mathcal{B}_L}(\tilde{\mathbf{y}}_i)/(\|\hat{\mathbf{E}}_{\mathcal{B}_L}(\tilde{\mathbf{y}}_i)\|_2 + \gamma), \forall i$

 Calculate the loss: $\mathcal{L}(\theta) = \frac{1}{|\mathcal{B}|} \sum_{i=1}^{|\mathcal{B}|} \|f_\theta(\tilde{\mathbf{y}}_i) - \mathbf{v}_{\mathcal{B}_L}(\tilde{\mathbf{y}}_i)\|_2^2$

 Update the model parameter: $\theta = \theta - \eta \nabla \mathcal{L}(\theta)$

end for

return f_θ

Algorithm 2 $\text{perturb}(\mathbf{x})$

 Sample the power $m \sim \mathcal{U}[0, M]$

 Sample the initial noise $(\epsilon_{\mathbf{x}}, \epsilon_z) \sim \mathcal{N}(0, \sigma^2 I_{(N+1) \times (N+1)})$

 Uniformly sample the vector from the unit ball $\mathbf{u} \sim \mathcal{U}(S_N(1))$

 Construct training point $\mathbf{y} = \mathbf{x} + \|\epsilon_{\mathbf{x}}\| (1 + \tau)^m \mathbf{u}$, $z = |\epsilon_z|(1 + \tau)^m$

return $\tilde{\mathbf{y}} = (\mathbf{y}, z)$

Backward ODE

Instead of solving backward ODE in time, solve the equivalent ODE in z :


$$\frac{d(x, z)}{dz} = \left(\frac{dx}{dt} \frac{dt}{dz}, 1 \right) = (v(\tilde{x})_x v(\tilde{x})_z^{-1}, 1)$$

Start from some $z \leq z_{\max}$, end at $z = 0$

Map hemisphere to $z = z_{\max}$ hyperplane

To allow processing batches, map the prior on the hemisphere $r = z_{\max}$ to the hyperplane

$$p_{\text{prior}}(x) = \frac{2z_{\max}}{S_N(1)(\|x\|_2^2 + z_{\max}^2)^{\frac{N+1}{2}}}$$



In practice: sample on hemisphere and project.

The actual Backward ODE for real

Add new time variable such that $dz = zdt'$


$$\frac{d(x, z)}{dt'} = (v(\tilde{x})_x v(\tilde{x})_z^{-1} z, z)$$

z now converges exponentially towards 0 in the backward ODE (actually makes solving the ODE faster)

IHDM



PFGM



Results

Table 1: CIFAR-10 sample quality (FID, Inception) and number of function evaluation (NFE).

	Invertible?	Inception \uparrow	FID \downarrow	NFE \downarrow
PixelCNN [36]	\times	4.60	65.9	1024
IGEBM [8]	\times	6.02	40.6	60
ViTGAN [24]	\times	9.30	6.66	1
StyleGAN2-ADA [17]	\times	9.83	2.92	1
StyleGAN2-ADA (cond.) [17]	\times	10.14	2.42	1
NCSN [31]	\times	8.87	25.32	1001
NCSNv2 [32]	\times	8.40	10.87	1161
DDPM [16]	\times	9.46	3.17	1000
NCSN++ VE-SDE [33]	\times	9.83	2.38	2000
NCSN++ deep VE-SDE [33]	\times	9.89	2.20	2000
Glow [19]	\checkmark	3.92	48.9	1
DDIM, T=50 [30]	\checkmark	-	4.67	50
DDIM, T=100 [30]	\checkmark	-	4.16	100
NCSN++ VE-ODE [33]	\checkmark	9.34	5.29	194
NCSN++ deep VE-ODE [33]	\checkmark	9.17	7.66	194
<i>DDPM++ backbone</i>				
VP-SDE [33]	\times	9.58	2.55	1000
sub-VP-SDE [33]	\times	9.56	2.61	1000

VP-ODE [33]	\checkmark	9.46	2.97	134
sub-VP-ODE [33]	\checkmark	9.30	3.16	146
PFGM (ours)	\checkmark	9.65	2.48	104
<i>DDPM++ deep backbone</i>				
VP-SDE [33]	\times	9.68	2.41	1000
sub-VP-SDE [33]	\times	9.57	2.41	1000

VP-ODE [33]	\checkmark	9.47	2.86	134
sub-VP-ODE [33]	\checkmark	9.40	3.05	146
PFGM (ours)	\checkmark	9.68	2.35	110

Results Summary

- Achieves best Inception and FID scores among normalizing flow models
- (10x ~ 20x) Faster than SDE models with similar architectures

Euler method with low step size




Figure 5: (a) Norm- $\sigma(t)$ relation during the backward sampling of VE-ODE (Euler). (b) Norm- $z(t')$ relation during the backward sampling of PFGM (Euler). The shaded areas mean the standard deviation of norms. (c) Number of steps versus FID score.

Miscellaneous

- Likelihood evaluation
- Latent representation

Limitations

- The mini-batch normalized field estimator is biased
- need large training batch to compensate

These are fixed in PFGM++

PFGM++

- Generalize to $z \in \mathbb{R}^D$
- equivalent to Diffusion models for $D \rightarrow \infty$
- Dispenses with the biased normalized field estimator
- No need to actually solve with the D additional dimensions. One is enough
- Transfer Hyperparameters from Diffusion Models to arbitrary D